

5 Modifying and Combining Data Sets

5.1 Stacking Data Sets Using rbind

The `rbind` function concatenates or stacks two or more data sets with all of the same variables but different observations. You might, for example, have data from two different locations or data taken at two separate times, but you need the data together for analysis.

You specify the new data set in the left hand side, then list the names of the old data sets you want to combine :

```
new-data-set = rbind(data-set-1,...,data-set-n)
```

The number of observations in the new data set will equal the sum of number of observations in the old data sets. The order of observations is determined by the order of the list of old data sets. All old data sets must contain the same set of variables and in the same order.

Example The following contains data for two entrances of an amusement park. The files has entrance (S or N), pass number, size of parties and their ages. The file for the north entrance has an N and the parking lot.

```
#South.dat
S 43 3 27
S 44 3 24
S 45 3 2
#North.dat
N 21 5 41 1
N 87 4 33 3
N 65 2 67 1
N 66 2 7 1
```

The following program reads the data and prints them. The third part combines the two data using `rbind`, it also creates a new variable `AmountPaid`, which tells how much each customer paid based on their age.

```
(southentrance <- read.table("./dataRaw/South.dat",
  head = FALSE, col.names = c("Entrance", "PassNumber", "PartySize", "Age"))

## Entrance PassNumber PartySize Age
## 1 S 43 3 27
## 2 S 44 3 24
## 3 S 45 3 2

(northentrance <- read.table("./dataRaw/North.dat",
  head = FALSE, col.names = c("Entrance", "PassNumber", "PartySize", "Age", "Lot")))

## Entrance PassNumber PartySize Age Lot
## 1 N 21 5 41 1
## 2 N 87 4 33 3
## 3 N 65 2 67 1
## 4 N 66 2 7 1

southentrance$Lot <- NA
both <- rbind(southentrance, northentrance)
both$AmountPaid <- c(0, 35, 27)[cut(both$Age, breaks=c(-99, 3, 65, 999), label = FALSE)]
```

The following are the results. Notice that the final data has missing values for the variable `Lot` for all observations which came from the south entrance.

```
both
## Entrance PassNumber PartySize Age Lot AmountPaid
## 1 S 43 3 27 NA 35
## 2 S 44 3 24 NA 35
```

## 3	S	45	3	2	NA	0
## 4	N	21	5	41	1	35
## 5	N	87	4	33	3	35
## 6	N	65	2	67	1	27
## 7	N	66	2	7	1	35

5.2 Interleaving Data Sets Using merge

The previous section explained how to stack data sets that have all the same variables at the same order, but different observations. However, you usually have data sets with similar set of variables, not necessarily in the same order. You could change the order and add in the missing variables before stacking the data sets. But a more general *join* operation would be more efficient, which performs all these operation for you. You need to use `merge` for such operations. Here is the general form:

```
new-data-set = merge(data-set-1, data-set-2, all=TRUE)
```

You specify the new data set in the left hand side, then names of two old data sets you want to combine, the `all=TRUE` denotes that all observations will be returned, whether they are from the first or the second data set. The number of observations in the new data set will equal the sum of number of observations in the old data sets. The order of observations is determined by the order of the list of old data sets. If one of the data sets has a variable not contained in the other data sets, values of that variable will be set to missing for observations from the other data sets.

Example We will use the amusement park data again.

```
#South.dat
S 43 3 27
S 44 3 24
S 45 3 2
#North.dat
N 21 5 41 1
N 87 4 33 3
N 65 2 67 1
N 66 2 7 1
```

Instead of stacking the two data sets, this program interleaves the data sets by pass number, entrance, party number and age. The data set `interleave` was created by combining the two data sets.

```
interleave <- merge(southentrance, northentrance[, c(2, 1, 3:5)], all = TRUE)
```

Here are the results. Notice that the results are the same even though the variables in the `northentrance` was permuted and missing values for `Lot` were automatically created.

```
interleave
## Entrance PassNumber PartySize Age Lot
## 1      S          43          3 27  NA
## 2      S          44          3 24  NA
## 3      S          45          3 2  NA
## 4      N          21          5 41  1
## 5      N          65          2 67  1
## 6      N          66          2 7   1
## 7      N          87          4 33  3
```

5.3 Combining Data Sets Using a One-to-One Match Merge

When you want to match observations from one data set with observations from another, use the `merge` function. If you know the two data sets are in *EXACTLY* the same order, you don't have to have any common variables between the data

sets. Typically, however, you will want to have, for matching purposes, a common variable or several variables which taken together uniquely identify each observation. This is important. Having a common variable to merge by ensures that the observations are properly matches. For example, to merge patient data with billing data, you would use the patient ID as a matching variable.

Merging data is a simple process. First name the new data set to hold the results, follow with a merge function, and list the data sets to be combined. Use the by argument to indicate the common variables:

```
new-data-set=merge(data-set-1,data-set-2,
  by=variable-list)
```

If you merge two data sets, and they have variables with the same names-besides the by variables, variables from both data sets will be kept, each renamed with suffix .x and .y.

Example A Belgian chocolatier keeps track of the chocolate sold each day. The code number for each chocolate and the number of pieces sold that day are kept in a file. In a separate file she keeps the names and descriptions of each chocolate as well as the code number. In order to print the day's sales along with the descriptions of the chocolates, the two files must be merged together using the code number. Here is a sample of the data.

```
## Sales dat
C865 15
K086 9
A536 21
S163 34
K014 1
A206 12
B713 29
## Description
A206 Mokka      Coffee buttercream in dark chocolate
A536 Walnoot    Walnut halves in bed of dark chocolate
B713 Frambozen  Raspberry marzipan covered in milk chocolate
C865 Vanille    Vanilla-flavored rolled in ground hazelnuts
K014 Kroon      Milk chocolate with a mint cream center
K086 Koning     Hazelnut paste in dark chocolate
M315 Pyramide   White with dark chocolate trimming
S163 Orbais     Chocolate cream in dark chocolate
```

```
descriptions <- read.fwf("./dataRaw/Chocolate.dat",
  widths = c(4, 10, 46),
  col.names = c("CodeNum", "Name", "Description"))
descriptions <- na.omit(descriptions)

sales <- read.table("./dataRaw/chocsales.dat", header = FALSE,
  col.names = c("CodeNum", "PiecesSold"))
```

The first parts of the program read the descriptions and sales data. The next part of the program creates a data set by merging the sales and the descriptions data set. The common variable CodeNum in the by statement is used for matching purposes.

```
merge(sales, descriptions, by = "CodeNum", all = TRUE)
```

##	CodeNum	PiecesSold	Name	Description
## 1	A206	12	Mokka	Coffee buttercream in dark chocolate
## 2	A536	21	Walnoot	Walnut halves in bed of dark chocolate
## 3	B713	29	Frambozen	Raspberry marzipan covered in milk chocolate
## 4	C865	15	Vanille	Vanilla-flavored rolled in ground hazelnuts
## 5	K014	1	Kroon	Milk chocolate with a mint cream center
## 6	K086	9	Koning	Hazelnut paste in dark chocolate
## 7	S163	34	Orbais	Chocolate cream in dark chocolate
## 8	M315	NA	Pyramide	White with dark chocolate trimming

The above output shows the data set after merging. Notice that the final data set has a missing value for `PiecesSold` in the eighth observation. This is because there were no sales for the chocolate. All observations from both data sets were included in the final data set whether they had a match or not.

5.4 Combining Data Sets Using a One-to-Many Match Merge

Sometimes you need to combine two data sets by matching one observation from one data set with more than one observation in another. Suppose you had data for every state in the U.S. and wanted to combine it with data from every county. This would be a one-to-many match merge, because each state observation matches with many county observations.

The statements for a one-to-many match merge are similar to the statements for a one-to-one match merge.

```
new-data-set=merge(data-set-1, data-set-2,
  by = variable-list, all.x = T/F, all.y = T/F)
```

The order of the data sets in the merge function matters. `all.x=T` and `all.y=F` mean that extra observations from `data-set-1` will be added to the results, even if there is no matching observation in `data-set-2` using the `by` variable. This is termed *left outer join*. `all.x=F` and `all.y=T` means the same for `data-set-2`, is termed *right outer join*. When both are "T", it is the default *full outer join* between two data frames.

You usually need a `by` argument for one-to-many merge. R uses the variables listed in the `by` argument to decide which observations belong together. Without any `by` variable for match, R simply produces a *Cartesian product* of the observations from two data sets. This may generate large data sets, slow and probably is not what you want.

If you merge two data sets, and they have variables with the same names-besides the `by` variables, variables from both data sets will be kept, each renamed with suffix `.x` and `.y`.

Example A distributor of athletic shoes is putting all its shoes on sale at 20 to 30% off the regular price. The distributor has two data files, one with information about each type of shoe and one with the discount factors. The first file contains one record for each shoe with values for style, type of exercise and regular price. The second file contains one record for each type of exercise and its discount. Here are the two raw data files:

```
## Shoes data
Max Flight      running 142.99
Zip Fit Leather walking 83.99
Zoom Airborne   running 112.99
Light Step      walking 73.99
Max Step Woven  walking 75.99
Zip Sneak       c-train 92.99
## Discount data
c-train .25
running .30
walking .20
```

To find the sale price, the following program combines the two data files:

```
regular <- read.fwf("./dataRaw/Shoe.dat", widths = c(15, 9, 6),
  col.names = c("Style", "ExerciseType", "RegularPrice"))
regular <- na.omit(regular)
discount <- read.table("./dataRaw/Disc.dat", header = FALSE,
  col.names=c("ExerciseType", "Adjustment"))
levels(regular$ExerciseType) <- levels(discount$ExerciseType)

prices <- merge(regular, discount, by="ExerciseType", all.x = TRUE)
prices$Newprice <- with(prices, round(RegularPrice * (1 - Adjustment)))
```

The first part reads the regular prices, creating a data set named `regular`. The second creates a data named `discount`. The levels of the `by` variable must match in order to perform the merge. The `merge` statement creates a data set named `prices`, merging the first two data sets, keeping all observations from `regular`, thus a *left outer join*. The output look like this:

##	ExerciseType	Style	RegularPrice	Adjustment	Newprice
## 1	c-train	Zip Sneak	93	0.25	70
## 2	running	Max Flight	143	0.30	100
## 3	running	Zoom Airborne	113	0.30	79
## 4	walking	Zip Fit Leather	84	0.20	67
## 5	walking	Light Step	74	0.20	59
## 6	walking	Max Step Woven	76	0.20	61

Notice that the values for Adjustment from the Discount data set are repeated for every observation in the regular data set with the same value of ExerciseType.

5.5 Merging Summary Statistics with the Original Data

Once in a while you need to combine summary statistics with your data, such as when you want to compare each observation to the group mean, or when you want to calculate a percentage using the group total. To do this, summarize your data using aggregate, and put the results in a new data set. Then merge the summarized data back with the original data using a one-to-many match merge.

Example 1 A distributor of athletic shoes is considering doing a special promotion for the top selling styles. The vice-president of marketing has asked you to produce a report. The report should be divided by type of exercise, and the percentage of sales for each style within its type. For each shoe, the raw data file contains the style name, type of exercise and the total sales for the last quarter.

Max Flight	running	1930
Zip Fit Leather	walking	2250
Zoom Airborne	running	4150
Light Step	walking	1130
Max Step Woven	walking	2230
Zip Sneak	c-train	1190

The program reads the raw data. Then it summarized the data with aggregate by the variable ExerciseType, creates a summary data named summarydata, containing a variable named Total, which equals the sum of the variable Sales.

```
shoes <- read.fwf("./dataRaw/Shoesales.dat",
                 widths = c(15, 9, 6), col.names = c("Style", "ExerciseType", "Sales"))
shoes <- na.omit(shoes)
summarydata <- aggregate(Sales ~ ExerciseType, data = shoes, FUN = sum)
names(summarydata)[2] <- "Total"
summarydata
```

##	ExerciseType	Total
## 1	c-train	1190
## 2	running	6080
## 3	walking	5610

In the second part of the program, the original data set shoes is merged with summarydata to make a new data set shoesummary. It also computes a new variable called Percent.

```
shoesummary <- merge(shoes, summarydata, by = "ExerciseType", all.x = TRUE)
shoesummary$Percent <- with(shoesummary, Sales * 100 / Total)
shoesummary
```

##	ExerciseType	Style	Sales	Total	Percent
## 1	c-train	Zip Sneak	1190	1190	100
## 2	running	Max Flight	1930	6080	32
## 3	running	Zoom Airborne	4150	6080	68
## 4	walking	Zip Fit Leather	2250	5610	40
## 5	walking	Light Step	1130	5610	20
## 6	walking	Max Step Woven	2230	5610	40

5.6 Adding Margins to a Data Set

The previous section discusses merging summary statistics with your data, that is summarizing variables across the observations. You may also need to summarize all variables corresponding to each observations, especially when all variables are recorded on the same scale. To do this, you will need to treat a data frame as a matrix and use `rowMeans` function to average over the columns instead of rows, put the results in a new data set or combine it with the original data sets.

Example A distributor of athletic shoes is evaluating the performance of her sales. The report should indicate the seasonal trend of the sales and the deviation of the sales from the overall mean. The raw data file contains the names of the sales and their sales in each season:

name	spring	summer	autumn	winter
Jane.Smith	14	18	11	12
Robert.Jones	17	18	10	13
Dick.Rogers	12	16	9	14
William.Edwards	15	14	11	10
Janet.Jones	11	17	11	16

The following program reads in the data and use `rowMeans` to calculate each people's average sales, its departure from the grand mean.

```
frame <- read.table("./dataRaw/sales.txt", header = TRUE, as.is = "name")
people <- rowMeans(frame[, 2:5])
(peopleeffect <- people - mean(people))

## [1] 0.30 1.05 -0.70 -0.95 0.30
```

The column means, which is the usual summary statistics for seasonal effects, are calculated in the similar way.

```
(season <- colMeans(frame[, 2:5]))

## spring summer autumn winter
##      14      17      10      13
```

Now we make a data frame that summarize the seasonal effects. We first create a data frame `newrow` to hold the results. The first column is the name, followed by the seasonal effects vectors.

```
newrow <- frame[1, ]
newrow[1,1] <- "Seasonal Effect"
newrow[1,2:5] <- season - mean(season)
```

The last task is to calculate the sales as deviation from the grand means and combine the seasonal effects with the original data, and add a column for sales performance.

```
newframe <- frame
newframe[,2:5] <- frame[,2:5] - mean(season)

data.frame(rbind(newframe, newrow),
            people=c(peopleeffect, mean(people)))

##           name spring summer autumn winter people
## 1 Jane.Smith   0.55   4.55  -2.5  -1.45   0.30
## 2 Robert.Jones 3.55   4.55  -3.5  -0.45   1.05
## 3 Dick.Rogers -1.45   2.55  -4.5   0.55  -0.70
## 4 William.Edwards 1.55   0.55  -2.5  -3.45  -0.95
## 5 Janet.Jones -2.45   3.55  -2.5   2.55   0.30
## 6 Seasonal Effect 0.35   3.15  -3.1  -0.45  13.45
```

Note that sales are highest in summer (+3.15) and lowest in autumn (−3.1). Robert Jones is the most effective sale person (1.05) and William Edwards is the least effective (−0.95).

5.7 Changing Observations to Variables Using reshape2

We have already seen ways to combine data sets, create new variables, and sort data. Now, using the reshape2 package, we will flip data, that is to change observations into variables.

The reshape2 transposes data sets, turning observations into variables or variables into observations. In most cases, to convert observations into variables, you first use the following statement.

```
long-data-set=melt(old-data-set, id-variables, measure-variables)
```

In the melt statement, old-data-set refers to the data set you want to transpose, and long-data-set refers to a long form of data suitable for transposition. id-variables are list of id variables. These variables are usually categorical, with combinations become the rows of the output data set, the "measure-variables" are variables that are stack onto each other in the long form, repeated for each unique combination of the id variables.

After melt, the dcast function can transpose the data:

```
new-data-set=dcast(long-data-set, by-variables ~ id-variables)
```

In the dcast statement, the by-variables are grouping variables separated by + that you want to keep as variables. These variables are included in the transposed data set, but they are not themselves transposed. The transposed data set will have one observation for each unique level of by-variables. The id-variables names the variable whose values become the variable names. The id-variables must occur only once in the data set within each by group. The function assumes that variables to transpose are in the field called value.

Example Suppose you have the following data about players for minor league baseball teams. You have the team name, player's number, the type of data (salary or batting average), and the entry:

```
Garlics 10 salary 43000
Peaches 8 salary 38000
Garlics 21 salary 51000
Peaches 10 salary 47500
Garlics 10 batavg .281
Peaches 8 batavg .252
Garlics 21 batavg .265
Peaches 10 batavg .301
```

You want to look at the relationship between batting average and salary. To do this, salary and batting average must be variables. The following program reads the raw data into a data frame. The data are first melted using all categorical variables, with the only one measurement variable Entry. The idea is to first make the data as long as possible, with each row denote one single variable corresponds to each unique combinations of id variables.

```
library(reshape2)
baseball <- read.table("./dataRaw/Transpos.dat",
                      head = FALSE, col.names = c("Team", "Player", "Type", "Entry"))
baseball.m <- melt(baseball,
                  idvars=c("Team", "Player", "Type"), measure.vars = "Entry")
head(baseball.m)

##      Team Player   Type variable  value
## 1 Garlics     10 salary      Entry 4.3e+04
## 2 Peaches      8 salary      Entry 3.8e+04
## 3 Garlics     21 salary      Entry 5.1e+04
## 4 Peaches     10 salary      Entry 4.8e+04
## 5 Garlics     10 batavg      Entry 2.8e-01
## 6 Peaches      8 batavg      Entry 2.5e-01
```

The data are then transposed using the dcast function. In the function, the by variables are Team and Player. You want those variables to remain in the output, and they define the new observations (you want only one observation for each team and player combination). The ID variable is Type, whose values (salary and batavg) will be the new variable names. By default, the variable to be transposed is value.

```
dcast(baseball.m, Team + Player ~ Type)
```

```
##      Team Player batavg salary
## 1 Garlics     10  0.28  43000
## 2 Garlics     21  0.26  51000
## 3 Peaches      8  0.25  38000
## 4 Peaches     10  0.30  47500
```

The wide format data set can be transposed back to long format, simply changing the by and id variables.

```
dcast(baseball.m, Team + Player + Type ~ variable)
```

```
##      Team Player  Type  Entry
## 1 Garlics     10 batavg 2.8e-01
## 2 Garlics     10 salary 4.3e+04
## 3 Garlics     21 batavg 2.6e-01
## 4 Garlics     21 salary 5.1e+04
## 5 Peaches      8 batavg 2.5e-01
## 6 Peaches      8 salary 3.8e+04
## 7 Peaches     10 batavg 3.0e-01
## 8 Peaches     10 salary 4.8e+04
```

5.8 Exercises

The following hypothetical data list the family income and spending for three years of three families.

```
#fam.txt
1 Jones
2 Smith
3 Jackson
#income.txt
1 96 40000 38000
1 97 40500 39000
1 98 41000 40000
2 96 45000 42000
2 97 45400 43000
2 98 45800 44000
3 96 75000 70000
3 97 76000 71000
3 98 77000 72000
```

Summarize the data by year for each family as a data frame:

	name	96_income	96_spend	97_income	97_spend	98_income	98_spend
1	Jackson	75000	70000	76000	71000	77000	72000
2	Jones	40000	38000	40500	39000	41000	40000
3	Smith	45000	42000	45400	43000	45800	44000

5.9 Exercises

Load swimming.txt file, and reshape it to get a data frame with the following columns (you may need to assign name to columns manually): treatment, subject, time, response.